

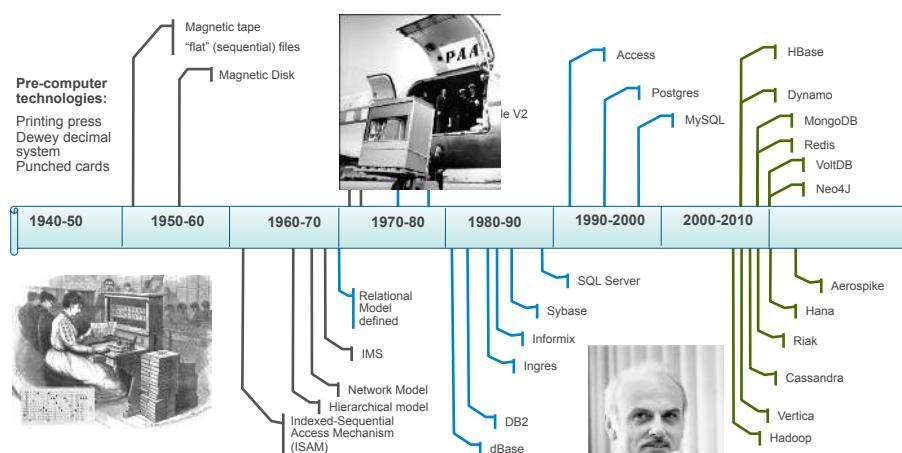
The slide features a dark background with various light gray icons related to databases and technology scattered across it. The title 'Next Generation Databases and Oracle' is prominently displayed in large white font. Below the title, the author's name 'Guy Harrison,' and title 'CTO dbKoda' are listed. To the right of the author's name is the dbKoda logo, which consists of a stylized 'd' and 'k' icon followed by the word 'dbKoda'. At the bottom left is a small dbKoda logo in a dark box.

The slide is divided into two main sections. On the left, there is a book cover for 'Next Generation Databases' by Guy Harrison, published by Apress. The book cover includes the author's bio: 'What every professional needs to know about the future of databases in a world of NoSQL and Big Data'. On the right, there is a graphic featuring several Star Wars characters (Yoda, Luke Skywalker, Han Solo, etc.) standing together. Overlaid on the graphic is the text 'YOU DON'T HAVE TO BE YODA TO USE DBKODA'. Below the book cover and the graphic are the author's contact details: Web: guyharrison.net, Email: guy.a.harrison@gmail.com, Twitter: [@guyharrison](https://twitter.com/guyharrison). There is also a logo for 'TOBA CAPITAL' featuring a triangle icon and the text 'TOBA CAPITAL'.

3 Database revolutions



History of databases

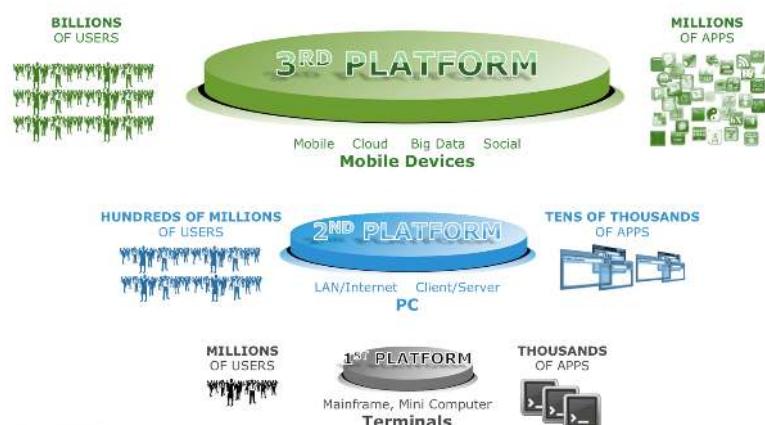


Why?

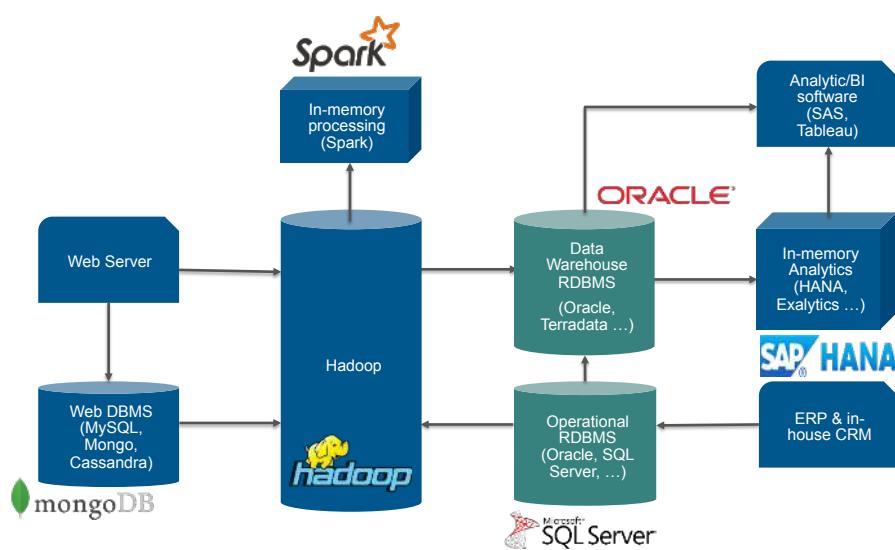
3rd Platform drives new demands on the database:

- Global high availability
- Data volumes
- Unstructured data
- Transaction rates
- Latency

A single architecture cannot meet all those demands.



It takes all sorts.



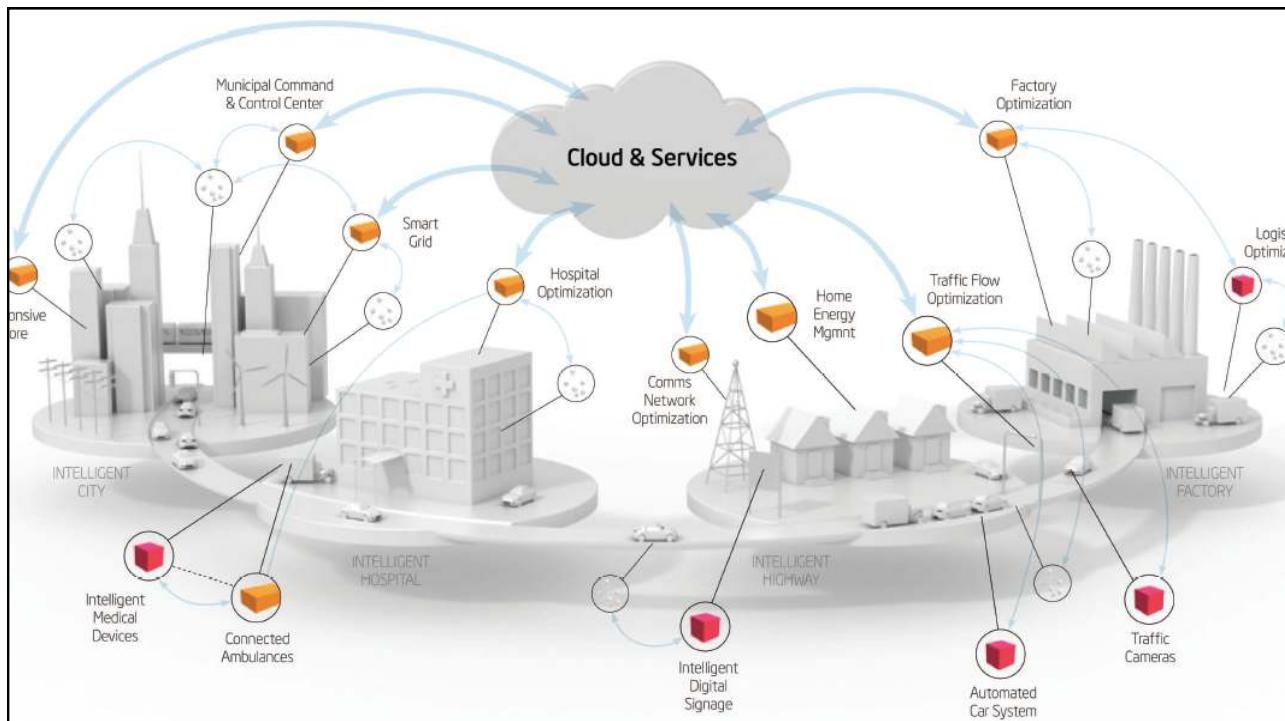
Big Data and Hadoop

dbKoda

The industrial revolution of data



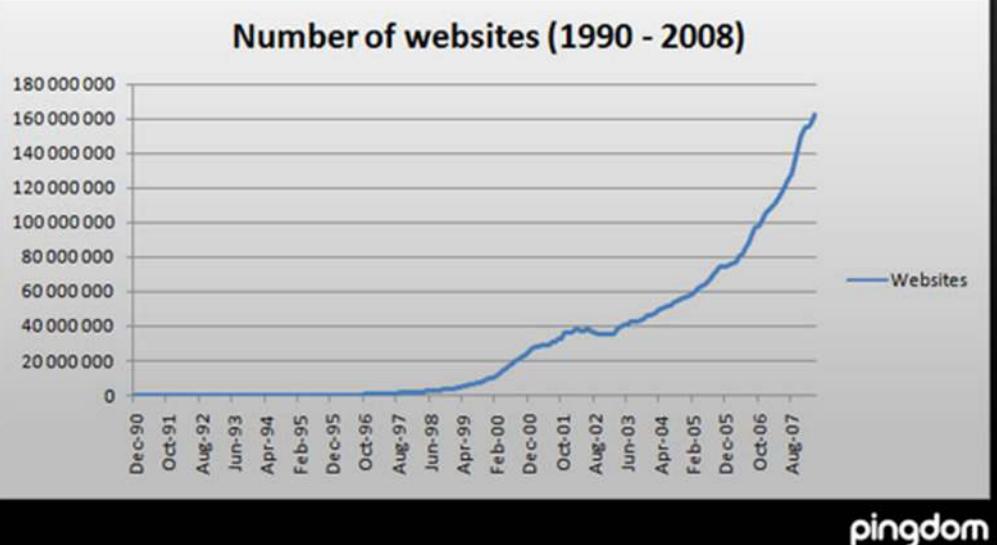


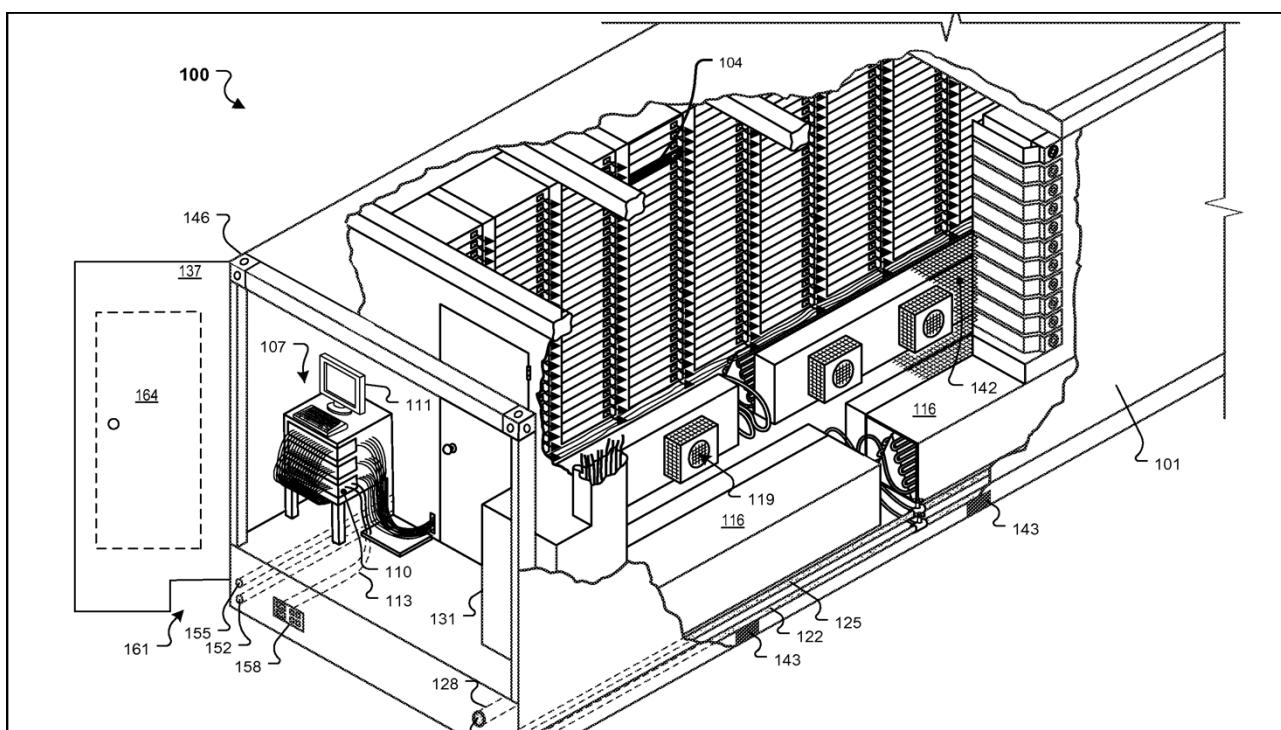
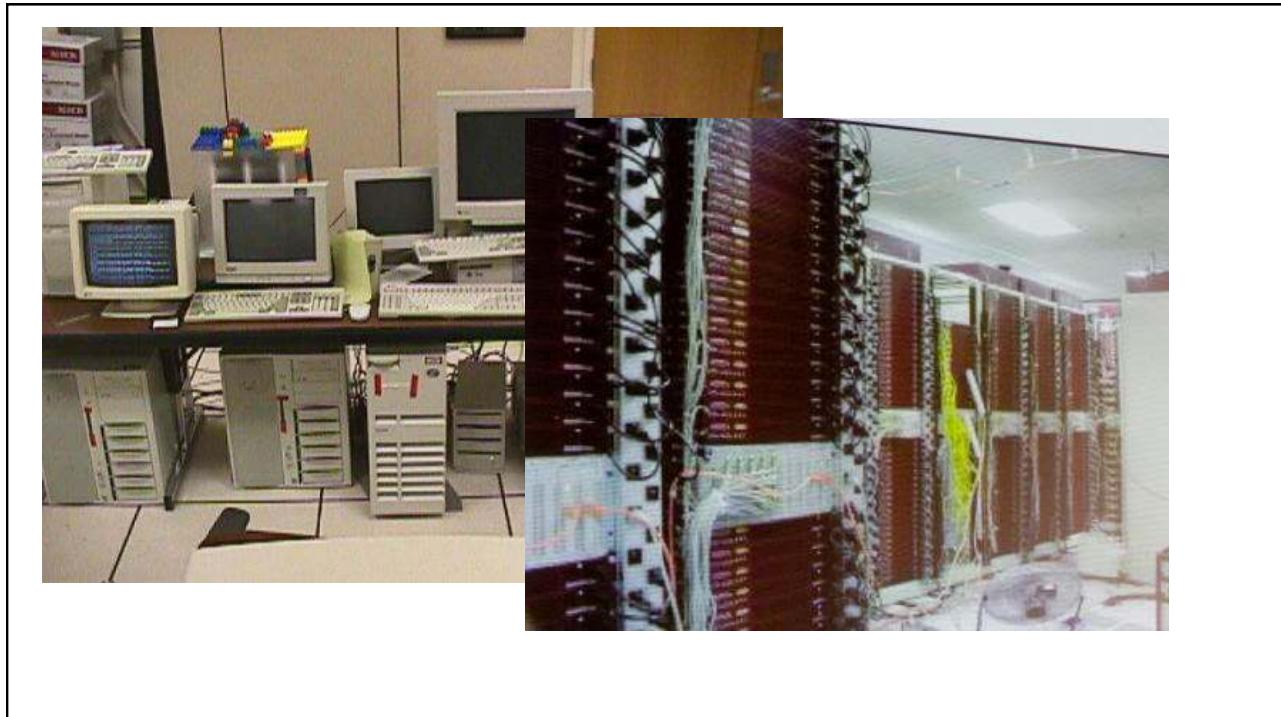


The instrumented human

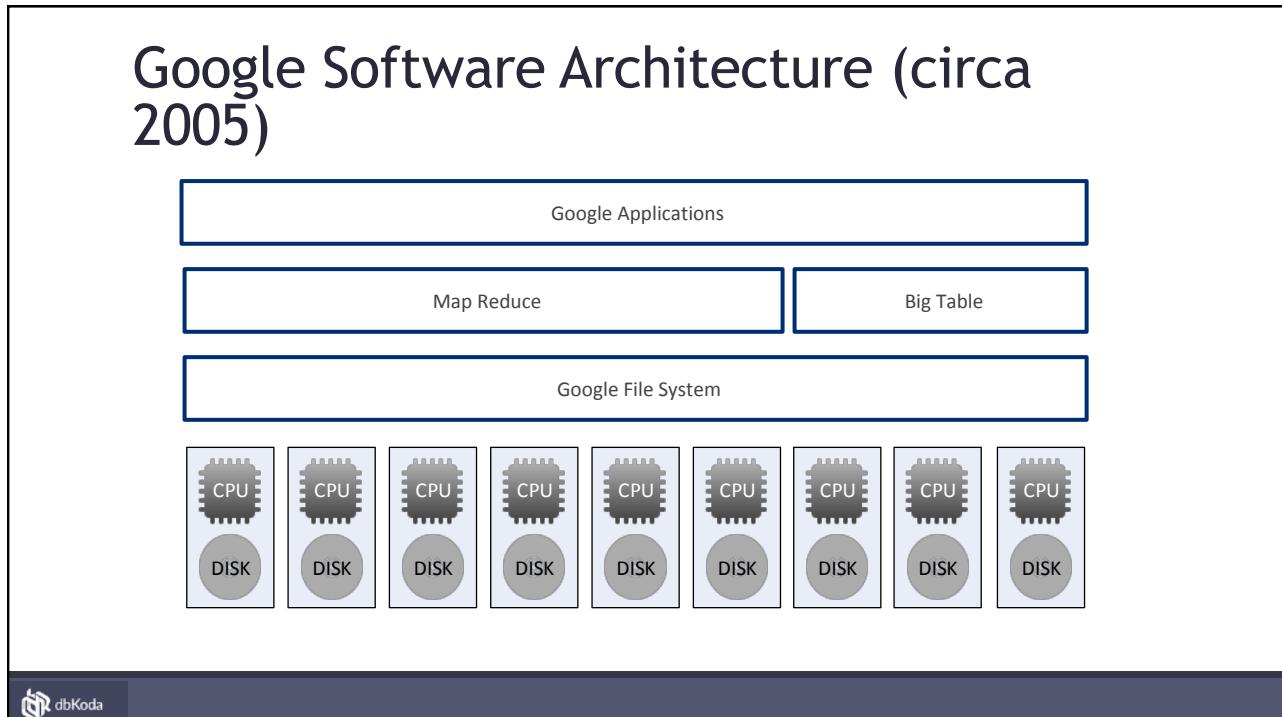
- Compass
 - Camera
 - Mike/earphones
 - Heads up display
 - Emotion/Attention monitor
- Bluetooth Personal Area Network
 - 3G/WiFi Wide Area Network
 - GPS
 - Storage
- Pulse, temp monitor
 - Silent alarms
 - Pedometer, sleep monitoring
-

Pioneers of big data









Hadoop at Yahoo

2010 (biggest cluster)

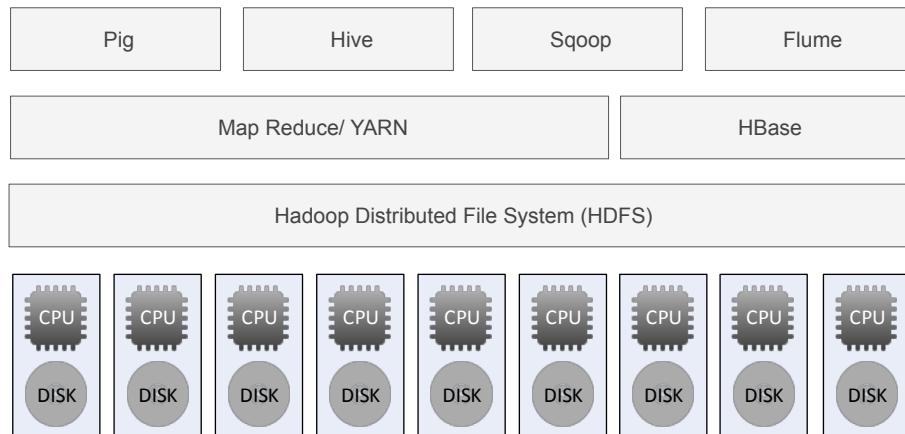
- 4000 nodes 16PB disk
- 64 TB of RAM
- 32,000 Cores

2014:

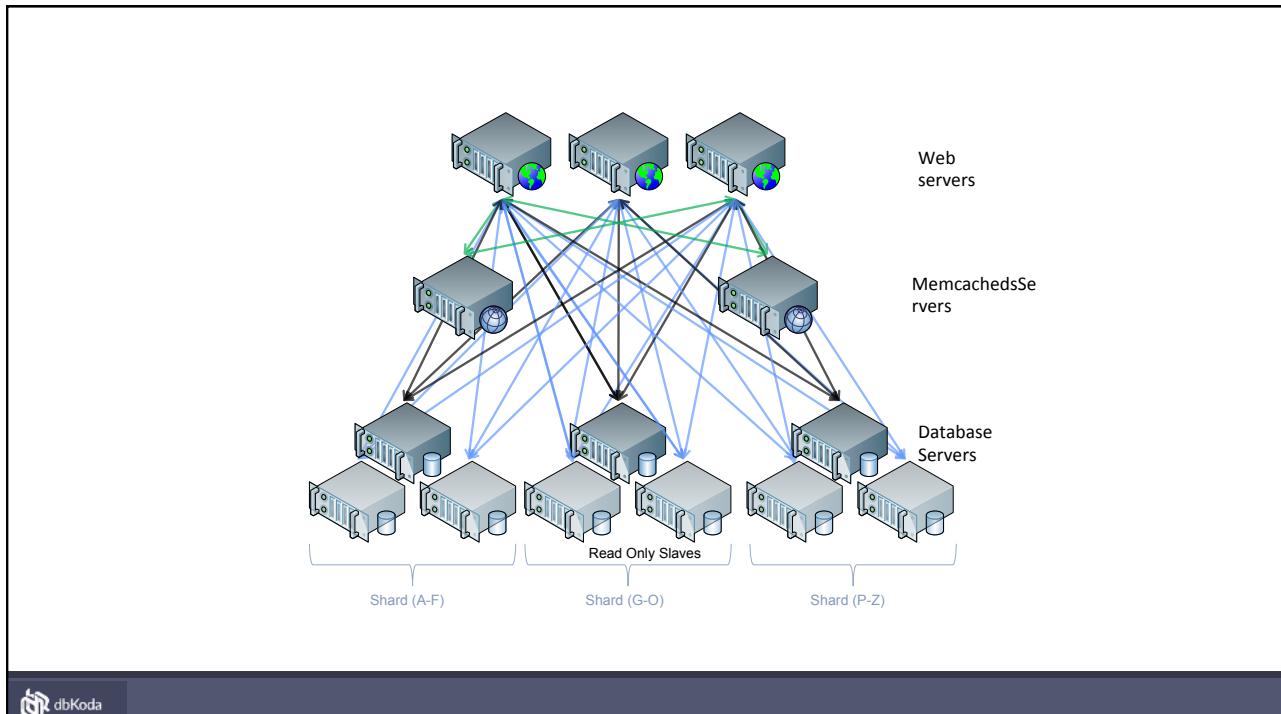
- 16 Clusters
- 32,500 nodes



Hadoop family



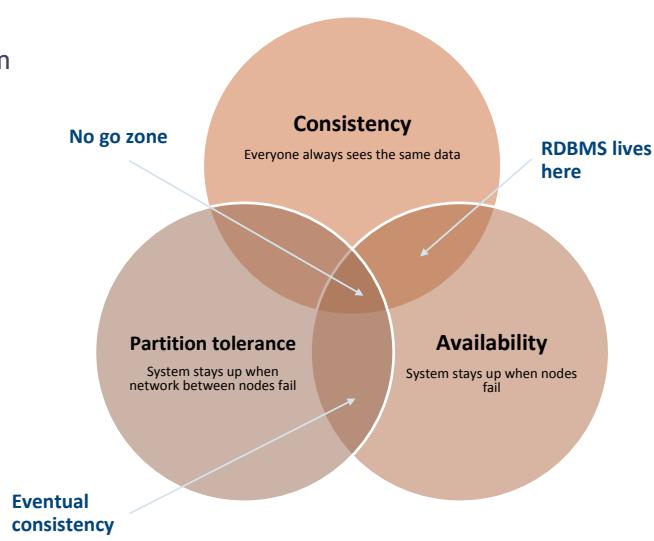
NoSQL



dbKoda

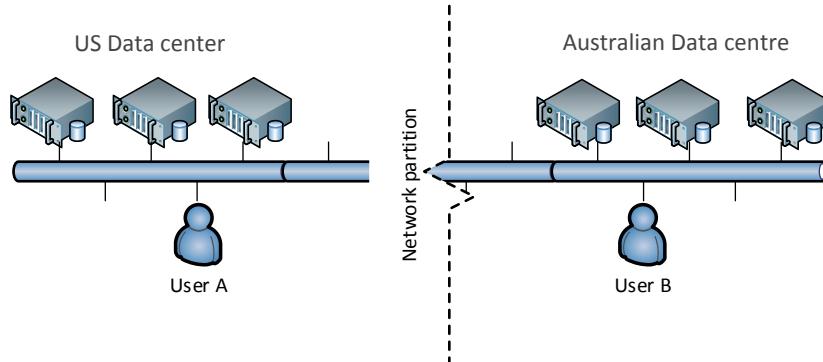
CAP Theorem says something has to give.

CAP (Brewer's) Theorem
says you can only have
two out of three of
Consistency, Partition
Tolerance, Availability



dbKoda

Network partition



Major influences on non-relational

Amazon Dynamo

- Eventually consistent transaction model
- Consistent hashing

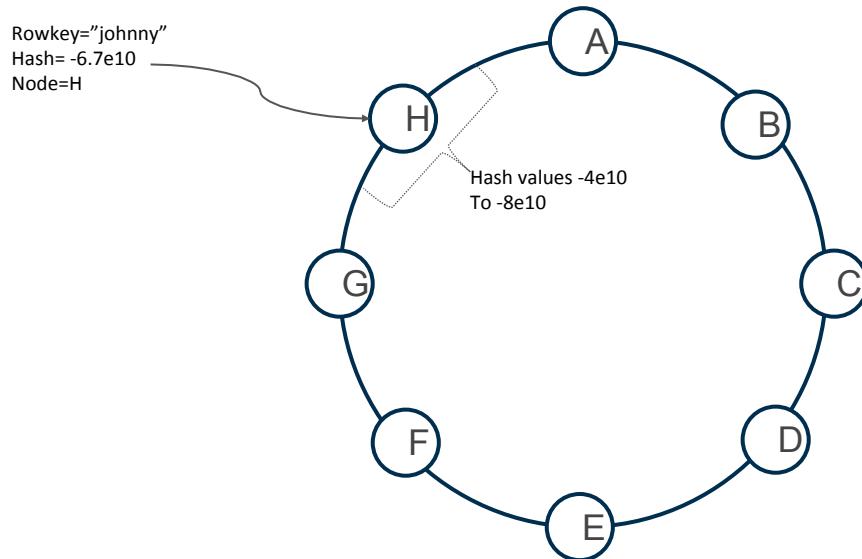
Google BigTable

- Column Family model for sparse distributed columnar data

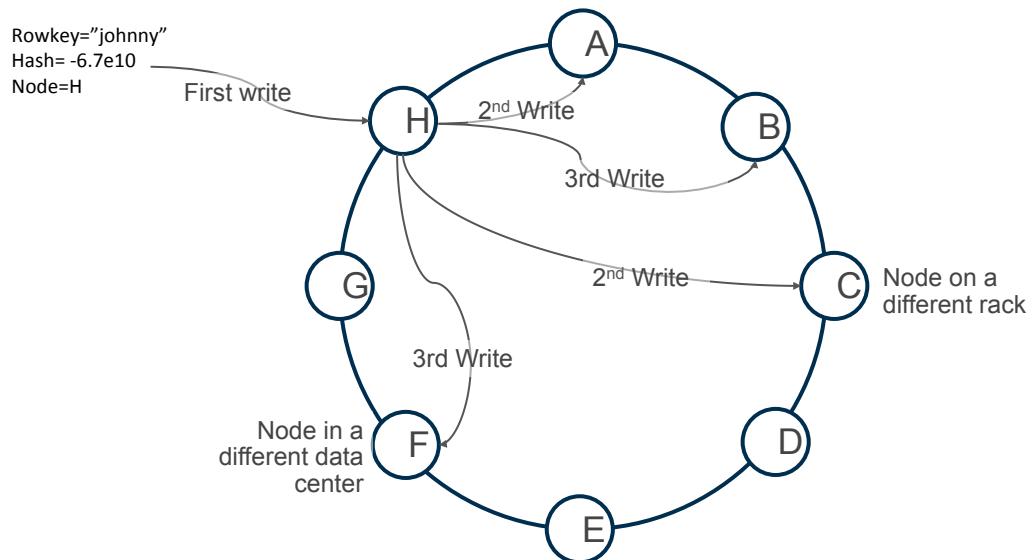
OODBMS and XML DBs

- Paved the way for the document database

Dynamo Consistent Hashing



Dynamo Consistent Hashing

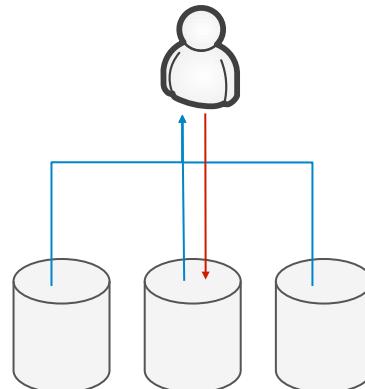


NWR - tunable consistency

N = No of copies of data

W = No of copies that must be written *synchronously* before returning control to program

R = No of copies to Read



$N=3 \ W=1 \ R=N$

Fastest write, slow but consistent reads

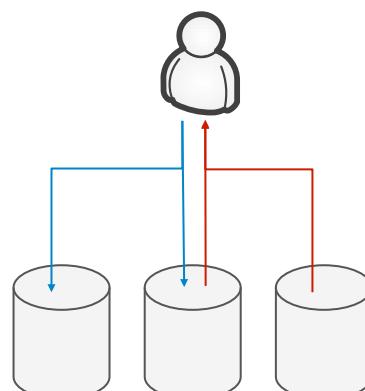
There will be 3 copies of the data.
A write request returns once the first copy is written – the other two can happen later.
A read request reads all copies to make sure it gets the latest version.
Data might be lost if a node fails before the second write.

NWR - tunable consistency

N = No of copies of data

W = No of copies that must be written *synchronously* before returning control to program

R = No of copies to Read



$N=3 \ W=2 \ R=2$

Faster writes, still consistent (quorum assembly)
There will be 3 copies of the data.

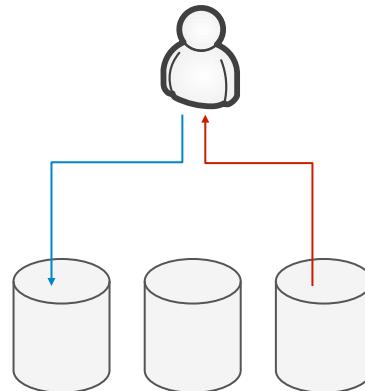
A write request returns when 2 copies are written – the other can happen later.
A read request reads two copies to make sure it has the latest version.

NWR - tunable consistency

N = No of copies of data

W= No of copies that must be written *synchronously* before returning control to program

R = No of copies to Read



N=3 W=1 R=1

Fast, but not consistent

There will be 3 copies of the data.

A write request returns once the first copy is written – the other two can happen later.

A read request reads a single version only: it might not get the latest copy.

Data might be lost if a node fails before the second write.

Revenge of the Object Nerds - Document databases

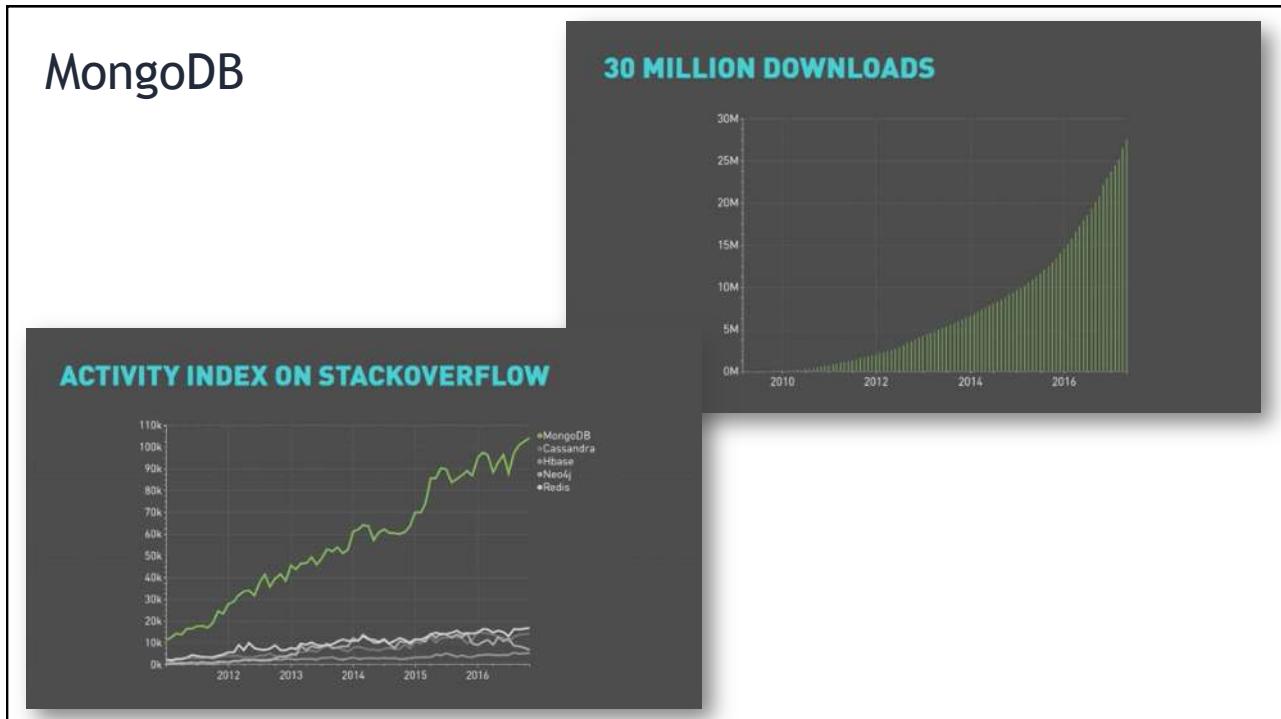
Structured documents - XML and JSON (JavaScript Object Notation) become more prevalent within applications

Web programmers start storing these in BLOBS in MySQL

Emergence of XML and JSON databases

Flexible schemas supported agile development and continuous deployment in a way RDBMS ALTER TABLE could not

<pre>* personnel.xml 4 P E 1 <?xml version="1.0" encoding="UTF-8"?> 2 <!DOCTYPE personnel SYSTEM "perso 3 <html><stylesheet type="text/css" 4 <personnel> 5 <person id="Big.Boss"> 6 <name> 7 <family>Boas</family> 8 <given>Big</given> 9 </name> 10 <email>chief@oxygennml.com 11 <link subordinates="One.W 12 </person> 13 <person id="One.worker"> 14 <name> 15 <family>Worker</family> 16 <given>One</given> 17 </name> 18 <email>one@oxygennml.com< 19 <link manager="Big.Boss"/> 20 </person> 21 <person id="Two.worker"> 22 <name> 23 <family>Worker</family> 24 <given>Two</given> 25 </name> 26 <email>two@oxygennml.com< 27 <link manager="Big.Boss"/> 28 </person> 29 <person id="Three.worker"> 30 <name></pre>	<pre>* personnel.json 4 P E 1 ["personnel": {"person": [2 { 3 "id": "Big.Boss", 4 "name": { 5 "family": "Boss", 6 "given": "Big" 7 }, 8 "email": "chief@oxygennml.com", 9 "link": {"subordinates": "One.W 10 }, 11 "One.worker": { 12 "name": { 13 "family": "Worker", 14 "given": "One" 15 }, 16 "email": "one@oxygennml.com", 17 "link": {"manager": "Big.B 18 }, 19 "Two.worker": { 20 "name": { 21 "family": "Worker", 22 "given": "Two" 23 }, 24 "email": "two@oxygennml.com", 25 "link": {"manager": "Big.B 26 }, 27 "Three.worker": { 28 "name": { 29 "family": "Worker", 30 "given": "Three" 31 }, 32 "email": "three@oxygennml.com", 33 "link": {"manager": "Big.B 34 } 35] 36 } 37]}</pre>
---	--



This screenshot shows the MongoDB Compass interface. On the left, the database browser displays a list of databases and collections. In the center, a modal window titled "Group By Query" is open, allowing users to define a query for the "SampleCollection" database and "Salika_film" collection. The modal includes fields for "Initial Filter conditions", "Columns To Group By" (set to "Category"), and "Columns To Aggregate". Below the modal, the main interface shows the MongoDB shell query input field and the results pane displaying the output of the aggregation query.

```

localhost:27016 - 1 (new1.js) ✘ localhost:27016 - 1 (Tree Action) ✘ localhost:27016 - 1 (agg.js) ✘
< db.getSiblingDB("SampleCollections").getCollection("Salika_film").aggregate([
  { $match:{ "Actors.Last_name": "CAGE", "Actors.First_name": "JOHN" },
  { $group:{ "_id": "$Category", "Count":{$sum:1} } },
  { $sort:{ "Count":-1 } }
])

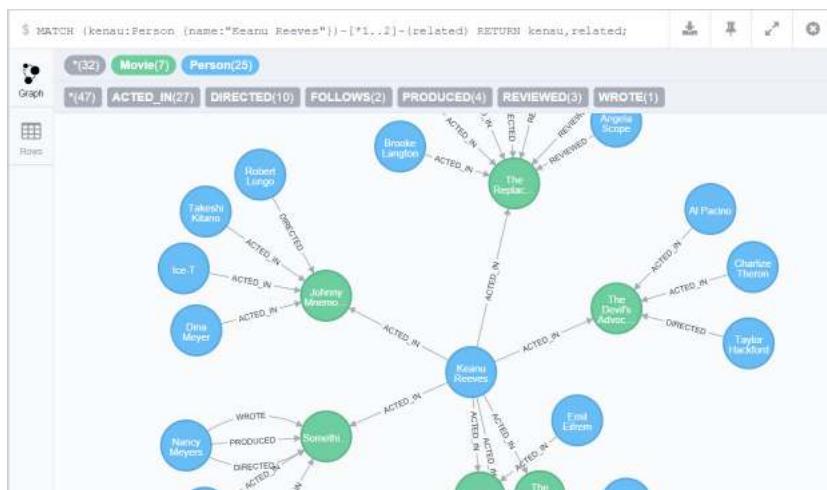
```

```

localhost:27016 - 1 (agg.js)
www.woot.com
+ db.getSiblingDB("SampleCollections").getCollection("Salika_film").aggregate([
  { $match:{ "Actors.Last_name": "CAGE", "Actors.First_name": "JOHN" },
  { $group:{ "_id": "$Category", "Count":{$sum:1} } },
  { $sort:{ "Count":-1 } }
])
+ JJJ
( { "id": { "Category": "Drama", "Count": 4 },
  { "id": { "Category": "Sci-Fi", "Count": 4 },
  { "id": { "Category": "Drama", "Count": 3 },
  { "id": { "Category": "Animation", "Count": 3 },
  { "id": { "Category": "Games", "Count": 2 },
  { "id": { "Category": "Comedy", "Count": 2 },
  { "id": { "Category": "Documentary", "Count": 2 }
}

```

Graph databases



dbKoda

No means yes!

Spark SQL

No
Only SQL

SHARK



APACHE DRILL

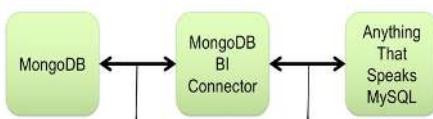


TERADATA

cloudera
IMPALA

presto

The MongoDB BI Connector: A “SQL Bridge”



Column-oriented Databases

Row orientation vs column orientation

ID	Name	DOB	Salary	Sales	Expenses
1001	Dick	21/12/60	67,000	78980	3244
1002	Jane	12/12/55	55,000	67840	2333
1003	Robert	17/02/80	22,000	67890	6436
1004	Dan	15/03/75	65,200	98770	2345
1005	Steven	11/11/81	76,000	43240	3214

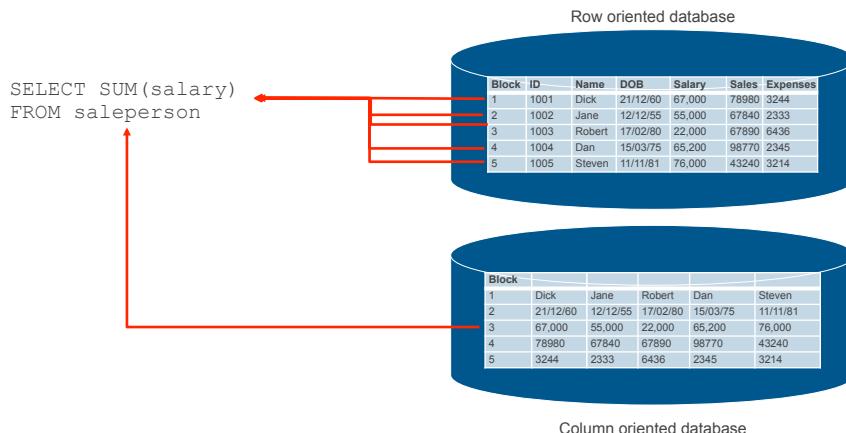
Row oriented database

Block	ID	Name	DOB	Salary	Sales	Expenses
1	1001	Dick	21/12/60	67,000	78980	3244
2	1002	Jane	12/12/55	55,000	67840	2333
3	1003	Robert	17/02/80	22,000	67890	6436
4	1004	Dan	15/03/75	65,200	98770	2345
5	1005	Steven	11/11/81	76,000	43240	3214

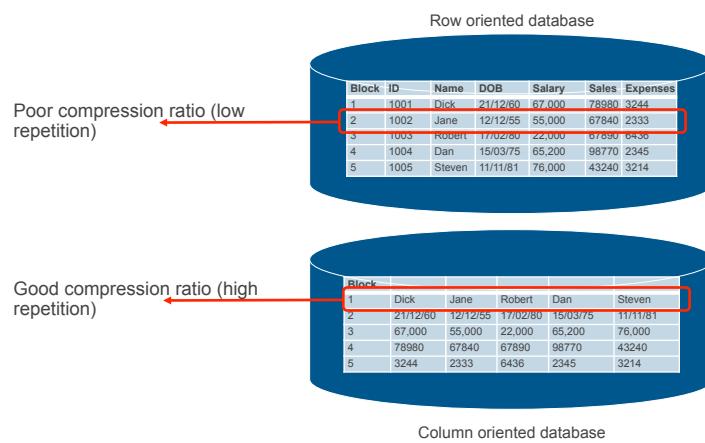
Column oriented database

Block	Dick	Jane	Robert	Dan	Steven
1	21/12/60	12/12/55	17/02/80	15/03/75	11/11/81
2	67,000	55,000	22,000	65,200	76,000
3	78980	67840	67890	98770	43240
4	3244	2333	6436	2345	3214

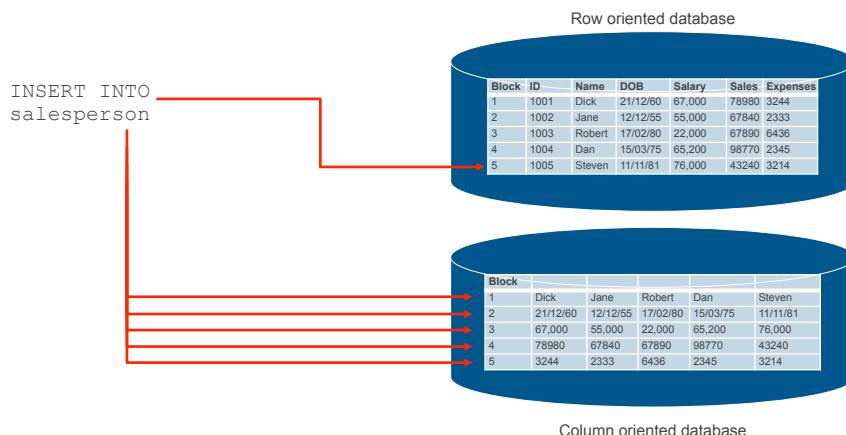
Analytical queries



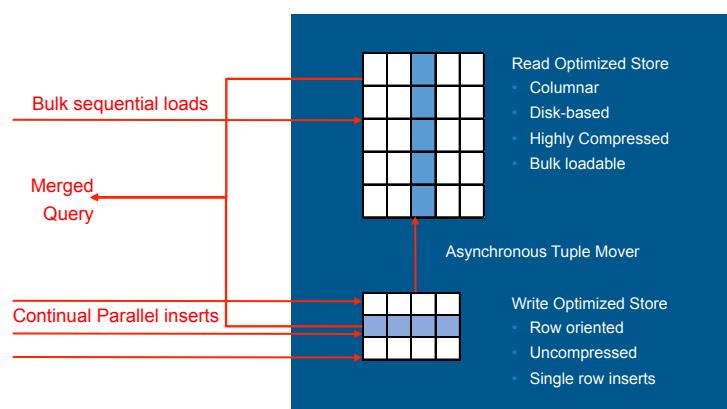
Compression



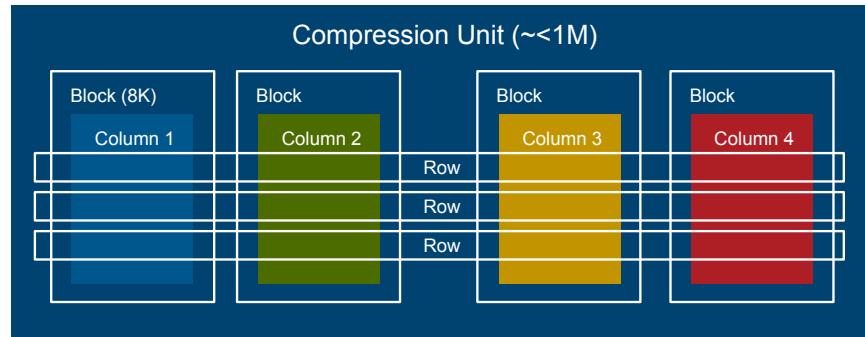
Inserts



C-Store (Vertica) solution for inserts



Exadata Hybrid Columnar Compression (EHCC)

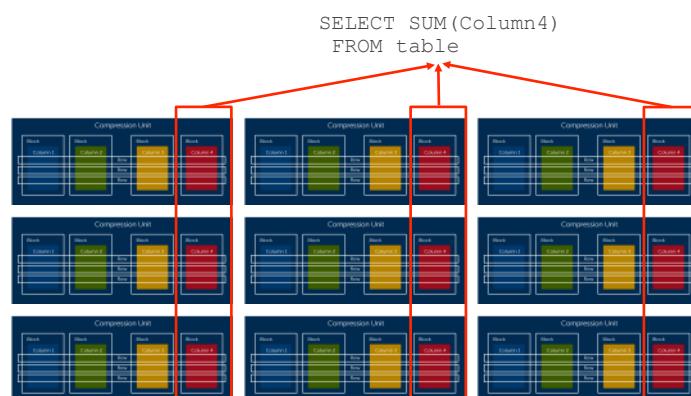


Exadata Hybrid Columnar Compression

Provides high compression ratio

Manageable impact on row read/write operations

Some optimization of analytic queries



SSD and in-memory databases

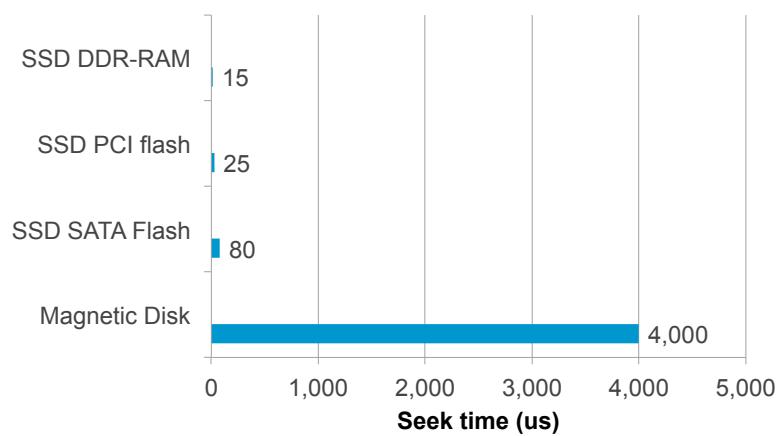
5MB HDD circa 1956



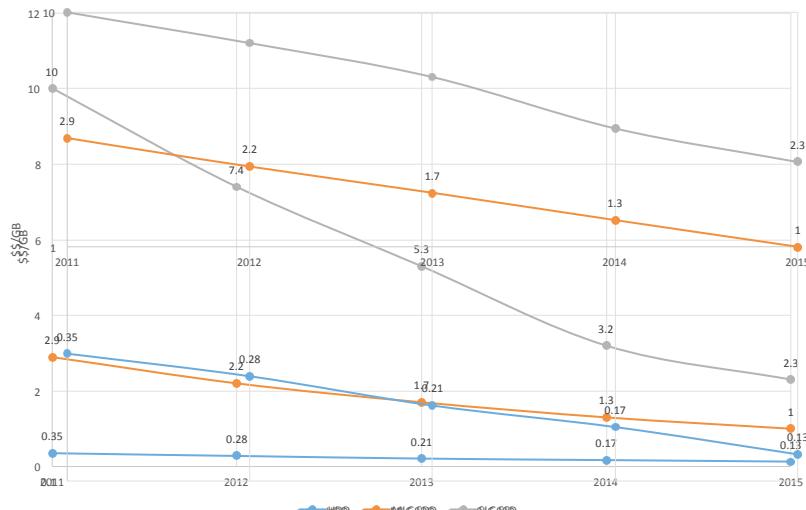
The more that things change....



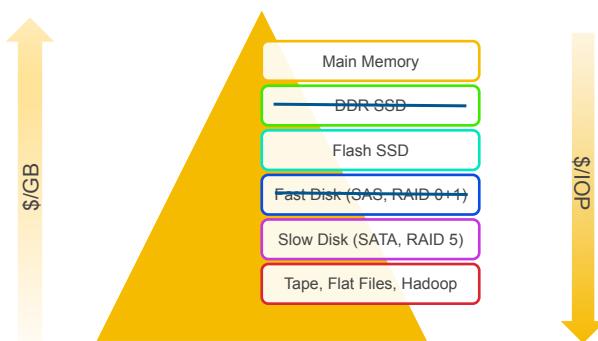
Cheaper by the IO



But not by the GB



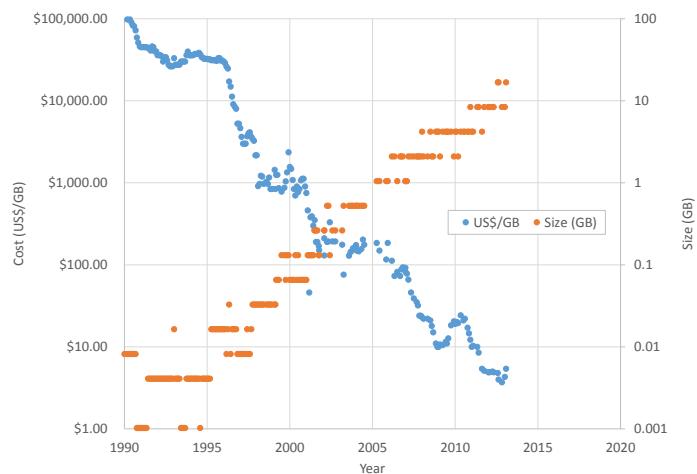
Tiered storage management



In-memory databases

Cost of RAM falling 50% each 18 months.

Some databases can fit entirely within the RAM of a single server or cluster of servers



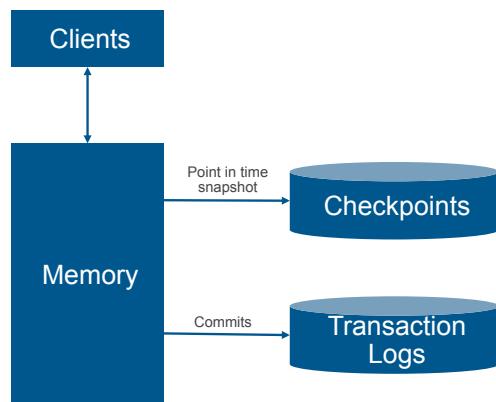
Oracle times ten

In-memory transactional database

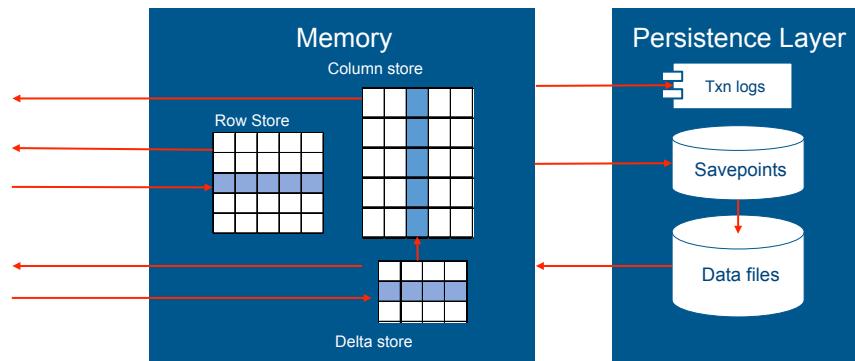
Disk-based Checkpoints and disk-based logging

By default, COMMITS are not durable (writes to the transaction log are asynchronous).

Can configure synchronous replication or synchronous log writes to avoid data loss



SAP HANA



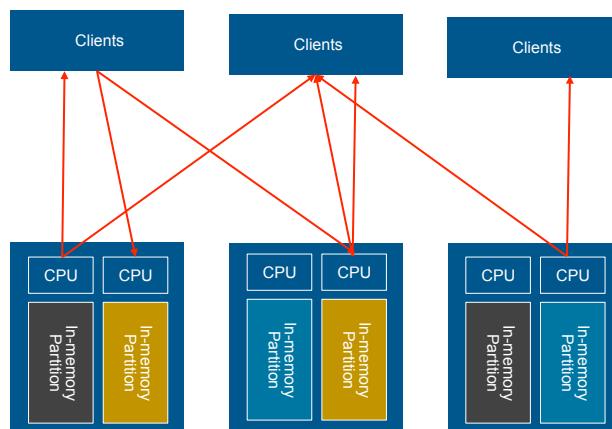
Note: Table must be either row or column – not both

VoltDB

Single threaded access to memory:
no latch/mutex waits

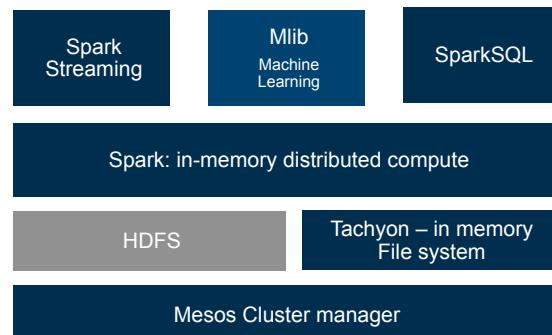
Transactions in self-contained stored procedures: minimal locking

K-Safety for COMMIT:
No sync waits

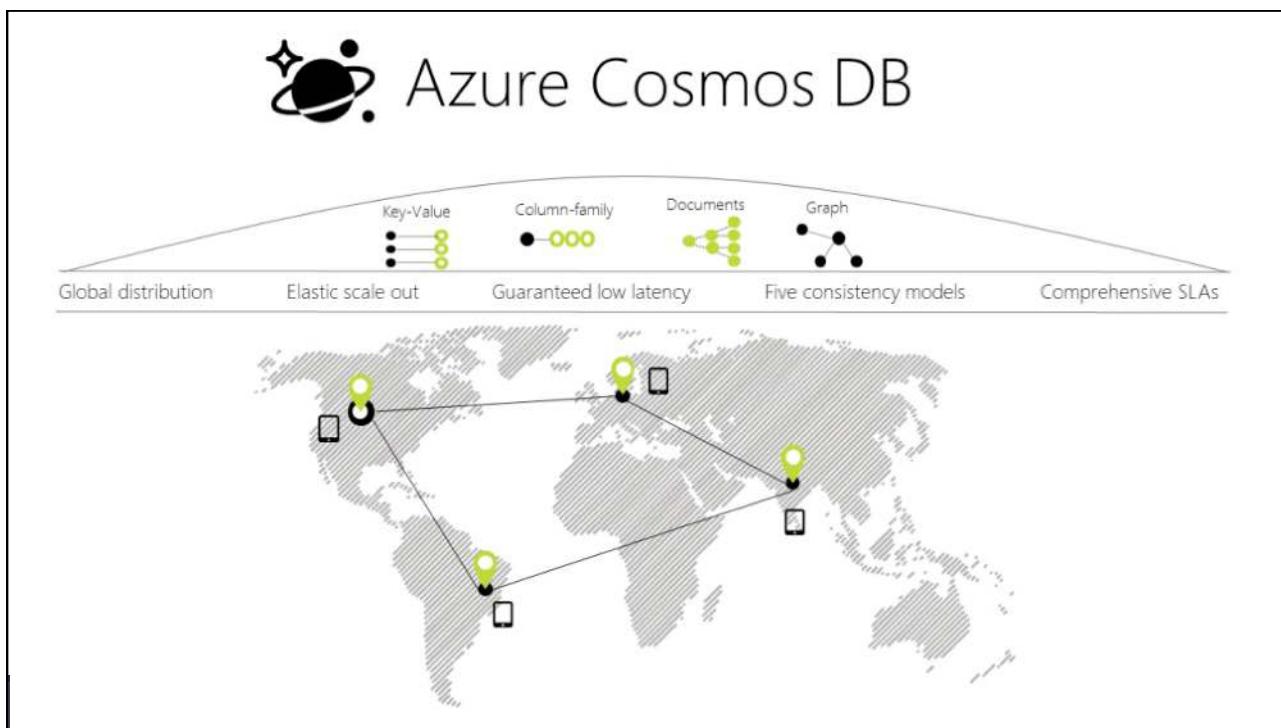
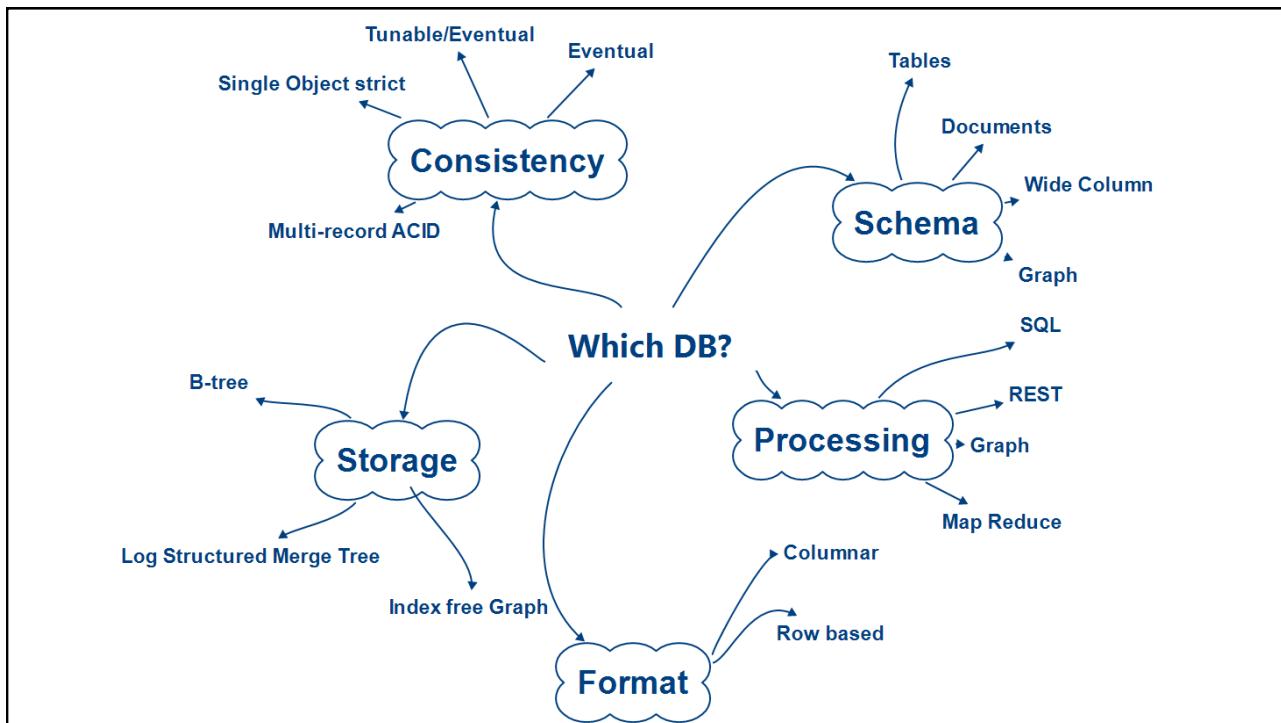


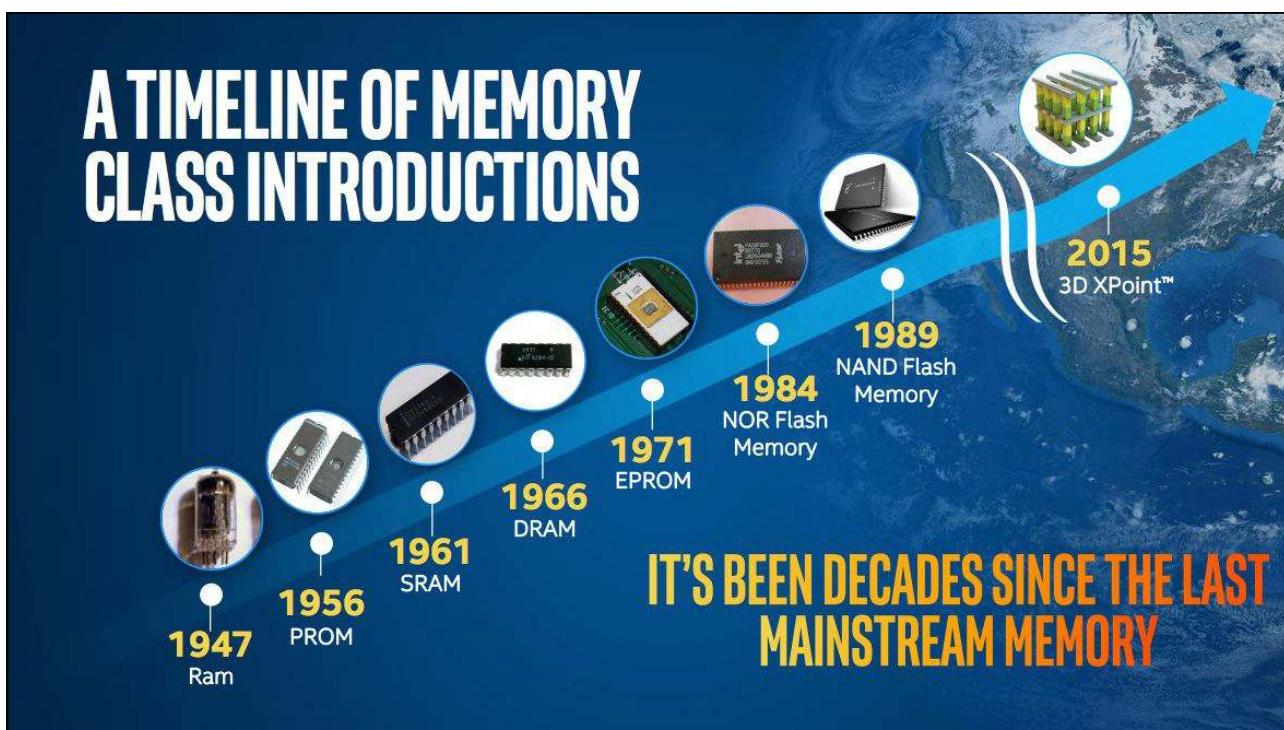
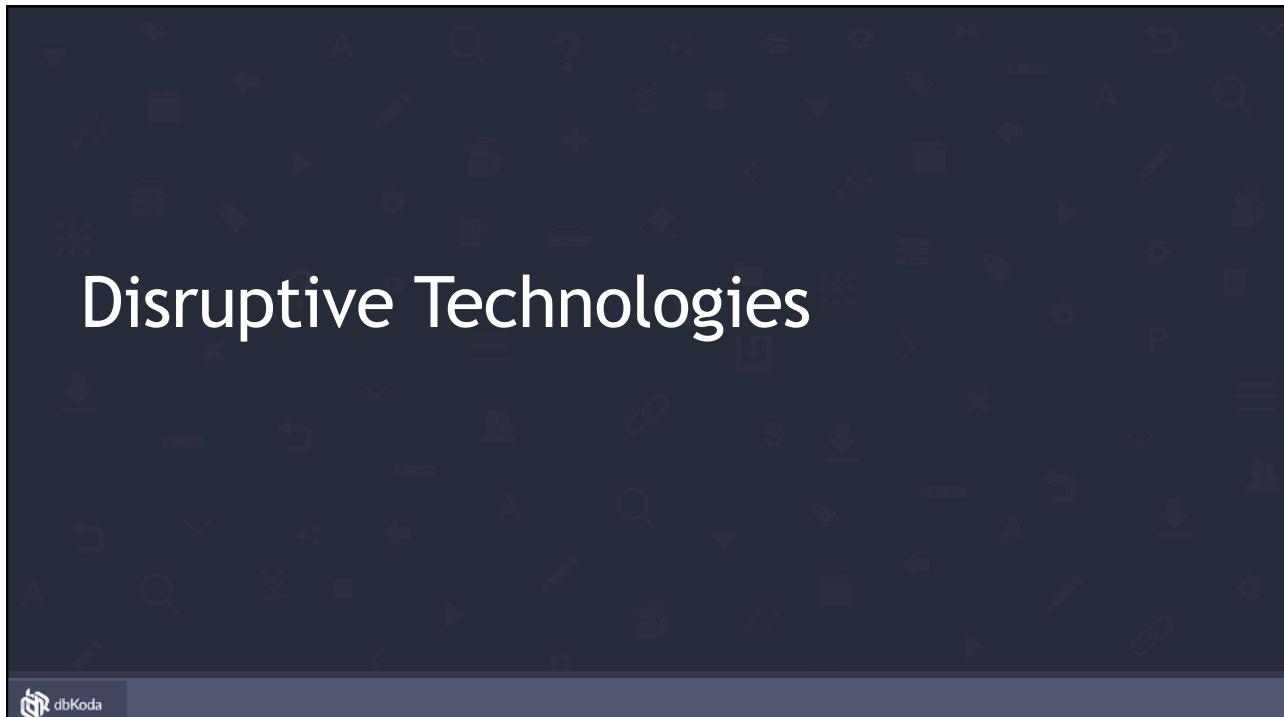
Spark (sort of) in-memory Hadoop

In Memory compute
HDFS compatible
Libraries for data processing,
machine learning, streaming,
SQL, etc
Python and Scala interfaces
Part of the Berkeley Data
Analytic Stack
Integrating into all Hadoop
distributions (and Cassandra)

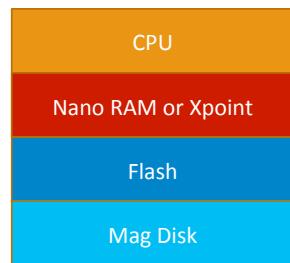


What will the database of
the future look like?





Universal Memory

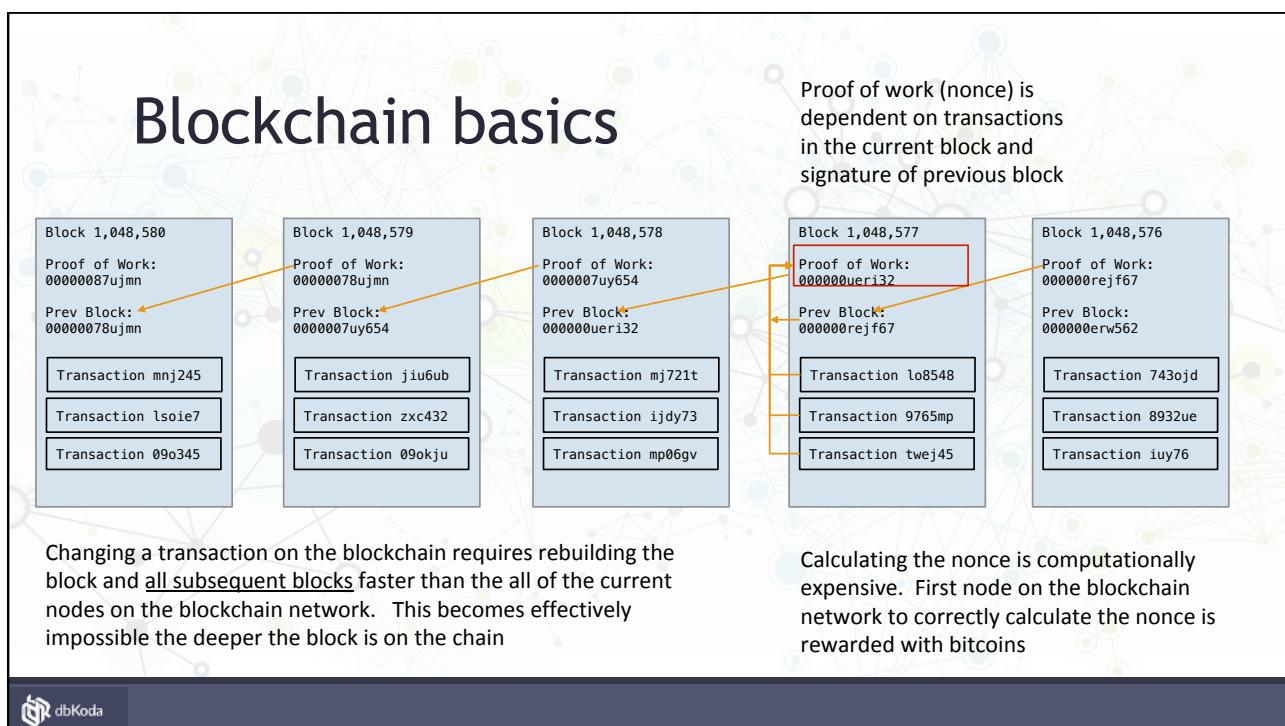


A single technology that provided performance & economics at all layers of storage would eliminate the rationale behind almost all existing database architectures



Blockchain





Blockchains and databases

The Blockchain represents a new type of database

- Public
- Distributed
- Tamper proof
- Immutable history

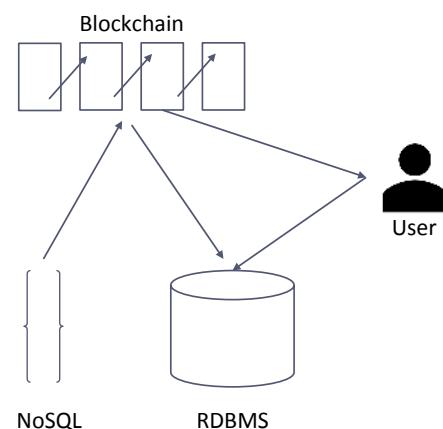
But doesn't provide traditional database features:

- Low throughput
- High latency
- No structure
- Limited volumes
- Excessively transparent

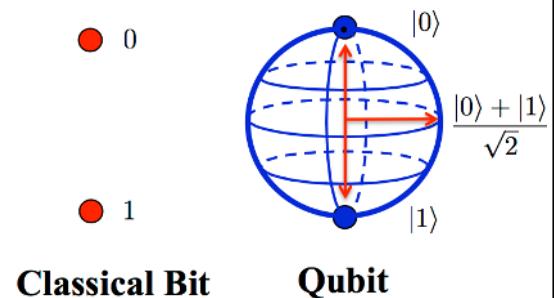


Databases of the future will leverage blockchain:

- Mediate changes to private data on public databases
 - Maintain the integrity of your personal data on a public database
- Facilitate transactions between disparate databases
 - Enable transactions between Oracle (shudder!) and MongoDB
- Assert proof of state
 - “Seal” data within a database



Quantum Computing



Thanks!

guy@southbanksofware.com
@guy_harrison
http://guy_harrison.net